



Module 3: Advanced Terraform Functions

Terraform-Intermediate-Day1-Module3

Learning Objectives

By the end of this module we will understand:

- Format strings for consistent naming
- Generate dynamic configuration files
- Combine multiple maps into one
- Build lookup tables from lists
- Flatten complex nested structures for use with *for_each*

Why Do We Need These Functions?

Imagine managing **multiple environments**:

Requirement	Example
Consistent Naming	web-01, web-02
Dynamic Config Files	Scripts change by environment
Combined Tags	Base + environment + resource
Flexible Security Rules	Multiple ports per rule group

The Limitation of Basic Interpolation

Basic interpolation works for simple strings:

```
"${var.project}-web"
```

But fails when we need:

- ✗ Zero-padded numbering (web-01)
- ✗ Loops in configuration files
- ✗ Combining multiple maps of tags
- ✗ Converting nested structures into flat lists

The Solution: Terraform Built-In Functions

Terraform includes powerful data transformation functions:

Function	Purpose
format()	Advanced string formatting
templatefile()	Render dynamic files
zipmap()	Build maps from lists
merge()	Combine maps
flatten()	Flatten nested lists

Quick Note: *path.module*

`path.module` gives the directory path of the current module.

```
templatefile("${path.module}/templates/config.tftpl", {...})
```

Format()

When You'll Use These in Real Projects

Scenario	Function Used
Naming EC2 instances	<code>format()</code>
Creating user data scripts	<code>templatefile()</code>
Mapping tiers to instance types	<code>zipmap()</code>
Combining tag layers	<code>merge()</code>
Breaking nested rules into individual rules	<code>flatten()</code>

Format()

Why Naming Standards Matter

Imagine a production environment:

Without Standard	With Standard
web1	web-01
web2	web-02
db1	db-01

Problems Without Formatting

- ✗ Hard to sort resources
- ✗ Hard to automate monitoring
- ✗ Hard to enforce conventions

Format()

Interpolation vs *format()*

Use Interpolation for Simple Joining

```
"${var.project}-${var.environment}-web"
```

Result:

```
myapp-dev-web
```

Use `format()` When we Need Special Formatting

```
format("%s-web-%02d", var.project, 1)
```

Result:

```
myapp-web-01
```

Format()

Understanding format() Syntax

```
format(format_string, value1, value2, ...)
```

The **format string** contains placeholders that get replaced with values.

Placeholder	Meaning
%s	String
%d	Decimal number
%02d	Zero-padded number (2 digits)

Format()

Breaking Down %02d

%02d means:

Symbol	Meaning
%	Start placeholder
0	Pad with zeros
2	Minimum width 2 digits
d	Decimal number

Examples

Code	Output
format("server-%02d", 1)	server-01
format("server-%02d", 9)	server-09
format("server-%02d", 12)	server-12 (already 2 digits, no padding needed)
format("server-%02d", 100)	server-100 (already 3 digits, no padding needed)

Format()

Real-World Naming Pattern

```
locals {
    name_prefix = "lab3-${var.project}-${var.environment}"
    formatted_server_name = format("%s-web-%02d", local.name_prefix, 1)
}
```

Result:

```
lab3-user1-dev-web-01
```

Format()

Breaking Down %02d

%02d means:

Symbol	Meaning
%	Start placeholder
0	Pad with zeros
2	Minimum width 2 digits
d	Decimal number

Examples

Code	Output
format("server-%02d", 1)	server-01
format("server-%02d", 9)	server-09
format("server-%02d", 12)	server-12 (already 2 digits, no padding needed)
format("server-%02d", 100)	server-100 (already 3 digits, no padding needed)

Templates

Why Do We Need Templates?

Imagine launching EC2 instances across environments:

Environment	Config Difference
Dev	Debug tools enabled
Staging	Monitoring enabled
Prod	Hardened security

Hardcoding scripts per environment =  Not scalable

Templates

What is templatefile()?

`templatefile()` reads a file and replaces placeholders with Terraform values.

```
templatefile(path, variables_map)
```

Parameter	Meaning
<code>path</code>	Path to template file
<code>variables_map</code>	Values to inject

Templates

Template File Basics

Template file example (`user_data.tftpl`):

```
#!/bin/bash
ENVIRONMENT="${environment}"
SERVER_NAME="${server_name}"
echo "Setting up $SERVER_NAME in $ENVIRONMENT"
```

Syntax	Purpose
<code>\$ {variable}</code>	Insert value

Templates

Template File Basics

Template file example (`user_data.tftpl`):

```
#!/bin/bash
ENVIRONMENT="${environment}"
SERVER_NAME="${server_name}"
echo "Setting up $SERVER_NAME in $ENVIRONMENT"
```

Syntax	Purpose
<code>\$ {variable}</code>	Insert value

Templates

Rendering the Template in Terraform

```
locals {  
  user_data = templatefile("${path.module}/templates/user_data.tftpl", {  
    environment = var.environment  
    server_name = "web-server"  
  })  
}
```

Result (environment = dev)

```
ENVIRONMENT="dev"  
SERVER_NAME="web-server"
```

Templates

Template Loops

Templates support loops for repeating sections.

Template:

```
#!/bin/bash
# Enable services
%{ for service in services ~}
systemctl enable ${service}
systemctl start ${service}
%{ endfor ~}
```

Output:

```
systemctl enable httpd
systemctl start httpd
systemctl enable sshd
systemctl start sshd
```

Terraform:

```
services = ["httpd", "sshd"]
```

Templates

Understanding ~} in Templates

- **The tilde (~}) removes extra newlines after loop blocks.**
- **Without ~}** → extra blank lines
- **With ~}** → clean formatting

Templates

Using *templatefile()* in Resources

```
resource "aws_instance" "web" {
    ami      = data.aws_ami.amazon_linux.id
    instance_type = "t3.nano"
    user_data  = local.user_data
}
```

zipmap()

The Problem: Parallel Lists

Imagine we define server tiers and their instance sizes separately:

```
tiers      = ["web", "db", "cache"]
server_types = ["t3.nano", "t3.micro", "t3.small"]
```

These lists are related by **position**, but Terraform cannot automatically link them.

Challenges:

- ✗ Hard to read
- ✗ Easy to mismatch values
- ✗ Not scalable

zipmap()

Creating Maps from Lists with zipmap()

The **zipmap()** function creates a map from two parallel lists - one list of keys and one list of values.

The Format

```
zipmap(keys_list, values_list)
```

Both lists must have the same length. The first item in the keys list pairs with the first item in the values list, and so on.

Basic Example

```
zipmap(["a", "b", "c"], [1, 2, 3])  
# Result: { a = 1, b = 2, c = 3 }
```

zipmap()

Step-by-Step Mapping

Position	Tier	Instance Type
1	web	t3.nano
2	db	t3.micro
3	cache	t3.small

Terraform Code

```
zipmap(var.tiers, var.server_types)
```

Output Example

```
{  
  web  = "t3.nano"  
  db   = "t3.micro"  
  cache = "t3.small"  
}
```

zipmap()

Practical Example: Instance Type Lookup

Imagine we have two related lists - server tiers and their corresponding instance types:

```
variable "tiers" {
    default = ["web", "db", "cache"]
}
variable "server_types" {
    default = ["t3.nano", "t3.micro", "t3.small"]
}
```

Use `zipmap()` to create a lookup table:

```
locals {
    tier_instance_types = zipmap(var.tiers,
        var.server_types)

    # Result: { web = "t3.nano", db = "t3.micro",
    #           cache = "t3.small" }
}
```

Now we can look up instance types by tier name:

```
resource "aws_instance" "web" {
    ami = data.aws_ami.amazon_linux.id
    instance_type =
        local.tier_instance_types["web"] # "t3.nano"
}
```

zipmap()

When to Use zipmap()

Use **zipmap()** when we have:

- we have **two related lists**
- Data is **positionally aligned**
- we need **name-based lookups**
- we want cleaner resource configuration

zipmap()

Important Rule: Lists Must Match

The lists must be the **same length**.

```
zipmap(["a","b"], [1,2,3]) ✘ Error
```

Why?

Terraform wouldn't know what value to assign to extra keys.

zipmap()

Real-World Use Cases

Scenario	Example
Tier to instance mapping	web → t3.nano
Region to AMI ID mapping	us-east-1 → ami-abc
Environment to CIDR mapping	dev → 10.0.1.0/24

merge()

The Real-World Problem

Imagine tagging strategy for cloud resources:

Layer	Example
Organization-wide tags	ManagedBy=Terraform
Environment tags	Environment=Dev
Resource-specific tags	Name=web-01

Without `merge()`:

- ✗ Tags scattered
- ✗ Duplicate definitions
- ✗ Hard to maintain consistency

merge()

Important Rule: Lists Must Match

merge() combines multiple maps into one.

```
merge(map1, map2, map3, ...)
```

If the same key appears multiple times → **last map wins**

merge()

Basic Example

```
merge(  
  { a = 1, b = 2 },  
  { b = 3, c = 4 }  
)
```

Result:

```
{ a = 1, b = 3, c = 4 }
```

merge()

Practical Example: Building Tags

```
variable "common_tags" {  
  default = {  
    ManagedBy = "Terraform"  
    Team      = "Platform"  
  }  
}
```

```
variable "environment_tags" {  
  default = {  
    CostCenter = "development"  
    Compliance = "standard"  
  }  
}
```

merge()

Layering Tags with **merge()**

```
locals {  
    all_tags = merge(  
        var.common_tags,  
        var.environment_tags,  
        {  
            Name      = local.server_name  
            Environment = var.environment  
        }  
    )  
}
```

Resulting Tags

ManagedBy = Terraform
Team = Platform
CostCenter = development
Compliance = standard
Name = lab3-user1-dev-web
Environment = dev

merge()

Override Behavior (Very Important)

```
merge(  
  { Environment = "default" },  
  { Environment = "dev" }  
)
```

Result:

```
Environment = dev
```

Rule

Later maps override earlier ones

merge()

Default + Override Pattern

```
locals {  
    default_tags = {  
        Environment = "unknown"  
        ManagedBy  = "Terraform"  
    }  
  
    final_tags = merge(  
        local.default_tags,  
        var.user_provided_tags  
    )  
}
```

merge()

Using `merge()` in Resources

```
resource "aws_instance" "web" {
  instance_type = "t3.nano"

  tags = merge(local.all_tags, {
    Name = "${local.server_name}-web"
  })
}
```

merge()

Real-World Use Cases

Scenario	Example
Tagging strategy	Org + Env + Resource tags
IAM policies	Base policy + team overrides
Security rules	Default rules + app-specific rules
Metadata	Global labels + custom labels

flatten()

The Problem: Nested Data

```
security_group_rules = [  
    {  
        name = "web"  
        ports = [80, 443]  
    },  
    {  
        name = "db"  
        ports = [3306, 5432]  
    }  
]
```

Easy for Humans

✓ Groups related ports

Hard for Terraform

✗ `for_each` cannot iterate nested lists directly

flatten()

What Does flatten() Do?

```
flatten([[1,2],[3,4]])
```

Result:

```
[1,2,3,4]
```

flatten() Transforming Rule Groups into Individual Rules

We want to convert:

web → [80,443]

db → [3306,5432]

Into:

web-80

web-443

db-3306

db-5432

flatten()

Nested for Expression Pattern

```
locals {
    flattened_rules = flatten([
        for rule in var.security_group_rules : [
            for port in rule.ports : {
                key = "${rule.name}-${port}"
                name = rule.name
                port = port
            }
        ]
    ])
}
```

flatten()

Resulting Flat List

```
[  
  { key = "web-80", name = "web", port = 80 },  
  { key = "web-443", name = "web", port = 443 },  
  { key = "db-3306", name = "db", port = 3306 },  
  { key = "db-5432", name = "db", port = 5432 }]  
]
```

flatten()

Converting List to Map for `for_each`

`for_each` requires a **map or set**, not a list.

```
locals {  
    ingress_rule_map = {  
        for rule in local.flattened_rules :  
            rule.key => rule  
    }  
}
```

Result:

```
{  
    "web-80" = { name="web", port=80 }  
    "web-443" = { name="web", port=443 }  
    "db-3306" = { name="db", port=3306 }  
}
```

flatten()

Using Flattened Data in a Resource

```
resource "aws_vpc_security_group_ingress_rule" "rules" {  
    for_each = local.ingress_rule_map  
  
    description      = each.value.name  
    from_port       = each.value.port  
    to_port         = each.value.port  
    ip_protocol     = "tcp"  
    cidr_ipv4      = "0.0.0.0/0"  
}
```

Key Takeaways

- **Use interpolation for simple string joining** ("\${var.a}-\${var.b}"), but use **format()** when we need special formatting like zero-padding
- **templatefile()** renders files with variable substitution and loops - perfect for user data scripts, configuration files, and any dynamic text content
- **zipmap()** creates a map from two parallel lists - useful for creating lookup tables from related data
- **merge()** combines maps with later values overriding earlier ones - the standard pattern for layering tags from multiple sources
- **Order matters in merge()** - put your defaults first and overrides last
- **flatten()** converts nested lists into flat lists - essential when your variable structure is hierarchical but we need individual items for `for_each`

Knowledge Check

Question 1

- What is the result of this expression?

```
format("server-%02d", 5)
```

- A)"server-5"
- B)"server-05"
- C)"server-005"
- D)"server-%02d5"

Solution:

B - The **%02d** format specifier means "decimal number, zero-padded to 2 digits". The number **5 becomes "05"**, so the full result is "server-05".

Knowledge Check

Question 2

- Given this nested structure:

```
rules = [  
    { name = "http", ports = [80] },  
    { name = "https", ports = [443] }  
]
```

Options:

- A) 1
- B) 2
- C) 3
- D) 4

After using the flatten pattern, how many items are in the resulting flat list?

Solution:

B - Each rule has one port, so flattening produces 2 items total: one for http port 80 and one for https port 443. The flatten pattern creates one item per port, not per rule group.

Knowledge Check

Question 3

- What is the result of this merge operation?

```
merge(  
  { Environment = "dev", Team = "alpha" },  
  { Environment = "staging", Owner = "bob" }  
)
```

- A) { Environment = "dev", Team = "alpha", Owner = "bob" }
- B) { Environment = "staging", Team = "alpha", Owner = "bob" }
- C) An error because Environment appears twice
- D) { Environment = ["dev", "staging"], Team = "alpha", Owner = "bob" }

Solution:

B - The `merge()` function combines maps, and when the same key appears multiple times, the last value wins. `Environment` appears in both maps, so the second map's value ("staging") overrides the first map's value ("dev"). The result includes all unique keys with the last-seen values.