



Terraform Lifecycle Configuration

Controlling How Resources Are Created, Updated, and Destroyed

Terraform-intermediate-Day1-Module-4

Terraform Lifecycle Configuration

Controlling How Resources Are Created, Updated, and Destroyed

Learning Objectives

By the end of this module, you will be able to:

 *Use create_before_destroy for zero-downtime updates* *Use prevent_destroy to protect critical resources* *Use ignore_changes to handle external modifications* *Decide when and how to combine lifecycle settings*

-  This module teaches how to make Terraform safer, smarter, and production-ready

⚠️ The Challenge

Why Do We Need Lifecycle Configuration?

Scenario: Managing Real Infrastructure

Your team manages:

- 🔒 Security Groups attached to running EC2 instances
- 💻 S3 buckets storing critical business data
- 🏷️ EC2 instances where AWS adds automatic tags

💡 The Problems

1 Replacement Order Issues

Terraform tries to **destroy** first, but AWS blocks deletion of attached security groups → Deployment fails

2 Accidental Deletion

Someone runs `terraform destroy` → Critical database or bucket deleted

3 Fighting External Tools

Security scanners add tags → Terraform tries to remove them every run

💡 We need **fine-grained control** over how Terraform handles resource lifecycle

❖ *The Solution*

The lifecycle Block

Terraform provides a special **meta-argument** called:

lifecycle

It lets you control **how Terraform treats a resource** during changes.

🔧 *What You Can Control*

Setting	Purpose
create_before_destroy	Avoid downtime during replacements
prevent_destroy	Protect critical resources
ignore_changes	Ignore externally modified attributes

- • The lifecycle block is placed **inside a resource block**
- • Works with **any Terraform resource**

 *Syntax*

Lifecycle Block Syntax

```
resource "aws_instance" "web" {  
    ami           = "ami-12345678"  
    instance_type = "t3.micro"  
  
    lifecycle {  
        create_before_destroy = true  
    }  
}
```

✓ Lifecycle rules are defined **inside the resource**

✓ You can combine multiple lifecycle settings

✓ Helps Terraform behave **safely in production**



create_before_destroy: Zero-Downtime Updates

Default Terraform Behavior (BAD for production)

01

Destroy old resource

02

Create new resource

✗ Causes downtime gap

With create_before_destroy = true

01

Create new resource

02

Switch dependencies

03

Destroy old resource

✓ No service interruption

```
lifecycle {  
  create_before_destroy = true  
}
```

🎯 Perfect for security groups, load balancers, IAM roles



Real-World Example

Real Example: Security Group Replacement

Without lifecycle control:

- ✗ Security group deleted
- ✗ EC2 instances temporarily lose access
- ✗ App downtime

With create_before_destroy:

- ✓ New security group created first
- ✓ Instances updated to use it
- ✓ Old group removed safely
- ✓ No traffic interruption

```
lifecycle {  
    create_before_destroy = true  
}
```

Best practice for **attached security groups**



The Setting

lifecycle { prevent_destroy = true }

What It Does

🚫 *Terraform refuses to destroy the resource*

🔴 *Stops the entire plan if destruction is attempted*

Examples:

- Production databases
- S3 buckets with business data
- Long-lived stateful systems

⭐ What Happens If Someone Runs Destroy?

If a resource has `prevent_destroy = true` and someone runs:

```
terraform destroy
```

Terraform responds with:

✗ *Error: Instance cannot be destroyed*

✗ *Terraform halts execution*

✗ *Nothing gets deleted*

This acts as a **safety brake** for production data



How to Remove a Protected Resource (Safely)

Sometimes deletion is necessary (e.g., system decommissioning).

Option 1 – Temporarily Comment It

```
lifecycle {  
  # prevent_destroy = true  
}
```

1. Run terraform apply
2. Then run terraform destroy

Option 2 – Set to False

```
lifecycle {  
  prevent_destroy = false  
}
```

Apply first → Then destroy

 ⚠ Always document and get approval before doing this



When Should You Use `prevent_destroy`?

Best candidates:



S3 buckets with important data



RDS databases



DynamoDB tables



Shared production infrastructure

 Note: This only protects from Terraform deletion, not manual AWS Console deletion

ignore_changes: Handling External Modifications

Sometimes attributes change **outside Terraform**:

- AWS adds system tags
- Security tools add compliance tags
- Auto Scaling modifies metadata

Terraform normally tries to **undo those changes** – causing unnecessary churn.

Solution

```
lifecycle {  
    ignore_changes = [attribute_name]  
}
```

Terraform will **ignore drift** for those attributes



Example: Ignoring External Tag Changes

```
resource "aws_instance" "web" {
    ami           = data.aws_ami.amazon_linux.id
    instance_type = "t3.micro"

    tags = {
        Name      = "web-server"
        Environment = "dev"
    }

    lifecycle {
        ignore_changes = [
            tags["LastScannedBy"],
            tags["ComplianceStatus"]
        ]
    }
}
```

✓ Terraform will not try to remove those tags

✓ Prevents conflict with security/compliance tools

🎯 *What Can Be Ignored?*

You can ignore:

01

Specific Attributes

```
ignore_changes = [instance_type, ami]
```

02

Nested Attributes

```
ignore_changes = [  
    tags["ExternalTag"],  
    root_block_device[0].volume_size  
]
```

03

Everything (Rare)

```
ignore_changes = all
```

⚠️ Using `all` means Terraform **won't update the resource after creation**

⚠️ *When Not to Use ignore_changes*

Do **not** use it to hide real problems.

Bad reasons:

- ✗ Avoiding proper configuration fixes
- ✗ Ignoring unexpected drift
- ✗ Skipping investigation of root cause

Good reasons:

- ✓ Known external automation
- ✓ AWS-managed attributes
- ✓ One-time initialization settings

🔗 Combining Lifecycle Settings

You can use multiple lifecycle rules together:

```
lifecycle {  
  create_before_destroy = true  
  ignore_changes = [  
    tags["LastScannedBy"],  
    tags["ComplianceStatus"]  
  ]  
}
```

But ⚠️ Important Rule

✖ prevent_destroy + create_before_destroy together doesn't make sense

If Terraform can't destroy, it also can't replace



Best Practices



Security Groups

```
lifecycle {  
  create_before_destroy = true  
}
```



Critical Data

```
lifecycle {  
  prevent_destroy = true  
}
```



External Tool Modifications

```
lifecycle {  
  ignore_changes =  
  [tags["LastScannedBy"]]  
}
```

- ❖ Use lifecycle deliberately – not by default



Always Document Lifecycle Usage

Future engineers (and you in 6 months) need context.

```
lifecycle {  
  # Create new SG before destroying old to avoid downtime  
  create_before_destroy = true  
  
  # Security scanner adds this tag automatically  
  ignore_changes = [tags["LastScannedBy"]]  
}
```

Good comments = fewer production surprises 

🎓 Key Takeaways

✓ *create_before_destroy*

Zero-downtime replacements

✓ *prevent_destroy*

Protection from accidental deletion

✓ *ignore_changes*

Avoid fighting external systems

✓ *Multiple settings*

Can be combined wisely

✓ *Lifecycle rules*

Make Terraform production-safe

Knowledge Check

Q1: Default Terraform replacement order?

- A) Create then destroy
- B) Destroy then create
- C) Update in place
- D) Prompt user

Solution: Q1: B